

When Speech Turns Violent

Exploring Threat and Violence in Incel Environments

Jithumol Kurian

Master's Degree Project in
Information Security
Spring term 2025
Supervisor: Lisa Kaati
Department of Computer
and Systems Sciences



**Stockholm
University**

Abstract

Introduction This study examines violent speech within blackpill.club, an online community for individuals who follow the Blackpill ideology, a more extreme form of the broader incel (involuntary celibate) movement. The Blackpill belief involves feelings of rejection, frustration, and anger toward women and society, often leading to a sense of hopelessness about relationships. Online violence, understood here as language that threatens or promotes physical harm to individuals or groups, is a growing concern in digital spaces. The Blackpill community also frequently expresses anger and resentment toward women and society, often using hostile or threatening rhetoric.

Research Question To address this issue, the study analyzes how often violent speech occurs on Blackpill.club, who the main targets are, and how severe the threats appear. Since the volume of online data makes manual annotation time-consuming, the study also compares the performance of two AI-based text annotation tools, ChatGPT and Grok, in detecting and classifying violent speech, alongside human annotation.

Method This study adopts a document survey as its research strategy and employs a mixed-method approach to analyze violent and threatening communication. Documentary research is used as the data collection method, where posts from the Blackpill.club forum are examined to identify patterns, frequency, and severity of violent speech.

Results The analysis of 1,182 comments from blackpill.club revealed that approximately 72% contained violent or threatening language, with women identified as the primary target group, followed by other categories including politicians, LGBTQ+, and various minority groups. Severity analysis of the violent comments indicated that approximately 36% were classified as high severity (explicit threats or encouragement of serious harm), 29% as moderate severity (less direct threats or harmful suggestions), and 34% as low severity (general hostile or aggressive language).

Discussion This study highlights the prevalence of violent and threatening speech on Blackpill.club, showing how such language targets specific groups and often uses violent content. While manual annotation offered the most reliable results, ChatGPT demonstrated stronger agreement with human annotators than Grok, though both tools showed limitations. Understanding these patterns is essential for developing better monitoring and intervention strategies to reduce online harm. These findings highlight the prevalence and intensity of violent speech in incel communities and demonstrate both the potential and the challenges of using AI to support large-scale detection of harmful online discourse.

Abbreviations

Abbreviation	Description
AI	Artificial Intelligence
NLP	Natural Language Processing
ML	Machine Learning
IoT	Internet of Things
GPU	Graphics Processing Unit
CPU	Central Processing Unit
API	Application Programming Interface

List of Figures

5.1	Distribution of violent and non-violent comments in the random sample of 600 comments.	14
5.2	Primary targets of violent speech	15
5.3	Severity Levels of Violent Comments (n=851)	16

List of Tables

4.1	Stages of the sampling and annotation process.	11
4.2	Severity Levels of Violent Comments	13
5.1	Krippendorffs alpha Agreement Scores Between Annotators and AI Tools	17
5.2	Inter-annotator Agreement (Krippendorffs alpha) for batch size 25	17

Contents

Abstract	iii
Abbreviations	iv
List of Figures	v
List of Tables	vi
Contents	vii
1 Introduction	1
1.1 Problem	2
1.2 Research Questions	3
1.3 Delimitations	3
2 Extended Background	4
2.1 Incels and social media	4
2.2 Understanding Violent Language in Incel Communities	5
2.3 Blackpill.club: An Online Hub for Blackpill Ideology	5
2.4 Understanding Dangerous Speech: A Framework for Identifying and Assessing Harmful Communication	6
3 Method	7
3.1 Data Collection	7
3.2 Data analysis	8
3.3 Ethical Considerations	8
4 Application Of Methodology	10
4.1 Data Collection	10
4.1.1 Data Extraction and Field Selection	10
4.1.2 Data Processing Workflow	10
4.1.3 Sampling Strategy	11
4.2 Data Analysis	11
4.2.1 Frequency of Violent Comments	12
4.2.2 Target Group Distribution	12
4.2.3 Severity of Violent Comments	12

5	Results	14
5.1	Frequency of Violent Speech	14
5.2	Primary Targets of Violent speech	15
5.3	Severity Levels of Violent speech	16
5.4	Comparing ChatGPT and Grok in Annotating Violent Speech	16
6	Discussion	18
6.1	RQ1: How frequently does violent or threatening speech appear on blackpill.club?	18
6.2	RQ2: Who are the primary targets of the violent speech?	19
6.3	RQ3: How severe is the violent speech?	19
6.4	RQ4: How do ChatGPT and Grok differ in annotating violent speech in black-pill.club, and which tool is more accurate and efficient	20
6.5	Limitations	20
6.6	Social and Ethical Implications	21
6.7	Conclusion	21
6.7.1	Future Research	22
6.8	Relevance to Cybersecurity	22
6.9	Use of AI tools	23
	Bibliography	26
	Appendices	27

1

Introduction

The internet has made it easier for people to connect, but it has also become a space where hateful and violent messages spread quickly. One group that has gained attention for its harmful communication is the “incel” community. The incel community is self-identified by a common inability to enter relationships or have sex with women [7]. The term “incel” was first coined by a woman named Alana, who initially created an online forum in the 1990s to discuss her struggles with loneliness and dating [23]. Originally, the incel community served as a space for mutual support, where individuals could share their experiences with isolation [5]. However, over time, many of these forums evolved into spaces of anger, resentment, and radicalization. As Jim Taylor notes in his BBC article, the movement shifted from a support group into a breeding ground for misogyny and extremist views, with some members expressing resentment not just toward women, but also toward men they perceive as more successful in romantic and social spheres [30].

To understand the incel community, it is crucial to first examine the broader online environment in which it exists, known as the manosphere [27]. The manosphere is a network of online communities that include a range of groups, each with distinct beliefs, but all united by a shared rejection of feminist ideals and gender equality. These communities arose from the Men’s Liberation (Men’s Lib) movement and now encompass subgroups such as Pick-Up Artists (PUA), Men’s Rights Activists (MRA), Men Going Their Own Way (MGTOW), and Involuntary Celibates (incels) [11]. While these groups may differ in their specific ideologies, they are linked by a common belief that modern gender dynamics, particularly feminist movements, are inherently oppressive to men. The manosphere, originally rooted in the desire for male equality, has, over time, evolved into an environment where many espouse negative, anti-feminist, and sometimes violent rhetoric [11].

A central concept that ties these groups together is the metaphor of the pills: blue pill, red pill, and black pill, which originates from the movie “The Matrix” [27]. The blue pill represents a state of ignorance, where individuals accept mainstream societal views on gender equality and feminism. In the manosphere, those who are “blue-pilled” are often criticized for not seeing what some believe is the way feminism harms men [11]. The red pill, in contrast, represents an awakening to the belief that men are oppressed by modern gender dynamics and that feminism has undermined traditional gender roles [17]. Those who adopt

a “red-pill” worldview believe that society unfairly favors women, positioning men as victims [28]. They claim that women are naturally drawn to only the most attractive, successful, and dominant men, often referred to as “Alpha males,” while rejecting those they consider less desirable, known as “Beta males” or “losers” [2]. The black pill, however, takes this belief even further, representing a more negative and hopeless view. Blackpillers believe that their situation is beyond change and that women are biologically programmed to choose “Alpha males” over them [28], no matter how much effort they put into improving themselves [17]. This belief leads to a sense of despair, where some even adopt the mindset of “Lie Down And Rot” (LDAR) [1], resigning themselves to their fate. Others turn this hopelessness into anger, sometimes using violent language and copying people like Elliot Rodger, whose actions have become a symbol for some in the community [1]. Blackpillers do not seek to improve their situation through traditional means, like the red-pill or blue-pill groups, but rather focus their energy on blaming women for their perceived failure and reinforcing a sense of rage against the world [3].

1.1 Problem

In recent years, the incel community has gained significant attention, particularly following high-profile violent incidents linked to individuals inspired by incel ideologies, such as the Isla Vista killings by Elliot Rodger and the Toronto van attack by Alek Minassian [27]. These incidents have raised alarms about the potential dangers of online radicalization, particularly in spaces that foster misogyny, extremism, and violence. The rapid growth of these online communities, especially those that embrace the blackpill ideology, poses a significant challenge to both social and digital security [5].

The amount of violent and threatening speech in online communities is increasing, making it important to understand how it works [25]. To address this problem, this research will investigate the frequency and severity of violent speech in blackpill.club, a forum where individuals who follow the blackpill ideology discuss their feelings of rejection and frustration, often expressing anger towards women and society [27]. The study will also examine who the targets are and what is the severity of the violent speech. Despite the growing concern about online violence, there is limited research specifically focusing on violent speech within the Incel forum, making this study particularly relevant. The existing literature tends to focus on incel communities broadly, but more detailed analyses are needed to understand the dynamics of violence in Incel forums.

Additionally, this study also addresses the growing challenge of analyzing large amount of data in online communities, where manual annotation is time consuming [19]. As online violence increases, AI models offer a more efficient way to process large amounts of data quickly. However, it is important to assess how these AI tools compare to human annotators in terms of accuracy and effectiveness. This study explores the use of AI tools, ChatGPT and Grok, to annotate violent speech in blackpill.club. By comparing the performance of ChatGPT and Grok to that of a human annotator using inter-rater reliability, this study aims to assess the strengths and limitations of these AI models in detecting violent language.

1.2 Research Questions

What are the characteristics, prevalence, and severity of violent and threatening communication on the incel forum blackpill.club, and how accurately and efficiently can the large language models ChatGPT and Grok identify and classify such content?

To address this overarching question, the study examines the following sub-questions:

- How frequently do violent speech appear in blackpill.club?
- Who are the primary targets of the violent speech?
- How severe is the violent speech?
- How do ChatGPT and Grok differ in annotating violent speech in blackpill.club, and which tool is more accurate and efficient?

1.3 Delimitations

This study focuses only on posts from the forum blackpill.club between 2021 and 2024. It examines violent and threatening language, not other types of harmful content such as insults, harassment, or hate speech without violent elements. The analysis is limited to three variables: violence classification, severity level, and target group. Only English language posts were included, and images, videos, or memes were excluded. The study also compares two AI tools ChatGPT and Grok and does not evaluate other language models. These choices define the scope of the study and ensure that the analysis stays focused on the research questions.

2

Extended Background

Social Media Platforms and Moderation Policies The role of social media platforms in shaping online discourse has evolved over time, with larger, more mainstream platforms now implementing stricter moderation policies [9]. These platforms are more likely to enforce guidelines to restrict harmful content, including hate speech and incitement to violence. However, some platforms, particularly smaller or niche forums, claim to offer a more permissive environment, often emphasizing “free speech” and allowing users to share content as long as it does not violate national laws [16]. This difference in moderation policies can have significant implications for the type and intensity of harmful content disseminated on these platforms. Studies by [10] indicate that the volume of extremist and antisemitic content on various platforms correlates with the level of moderation in place. On platforms with stricter moderation, such content is less prevalent, whereas on platforms with minimal oversight, the content tends to be more extreme and pervasive [10].

2.1 Incels and social media

The identification of incels in social networks has been explored in several studies. Papadamou et al. [20] conducted an analysis of YouTube content, revealing a substantial rise in inappropriate content generated by incel users. The study also observed that platforms such as YouTube often recommend this material to users [9]. The study suggests that social media platforms should implement algorithms to detect and prevent the spread of incel-related content to avoid further radicalization. In [6], the author discussed how the incel community has become a platform for young men to express anger towards women who rejected them. The study called for legal changes in communication laws to address these issues. While forensic activities may help reduce such harmful behaviors, the study emphasized the need for more methods to detect such problems in online communities [6].

investigated online responses to the 2018 Toronto attack by analyzing communication on incels.me and Reddit.com, using both qualitative and quantitative methods[13]. They found that after the incels.me forum was suspended, users migrated to Reddit. Rouda et al. examined why young men are attracted to incel communities [27]. They highlighted factors

like loneliness and anger towards women, but their study focused on terrorist acts rather than the detection.

2.2 Understanding Violent Language in Incel Communities

According to [16] violent and threatening refers to any form of written communication that explicitly or implicitly threatens physical harm or violence against individuals, groups, or entities. Research on violent language in Incel communities has gained significant attention within computational social science [15]. Various studies have employed natural language processing (NLP) techniques to analyze the language used in Incel forums. These studies focus on uncovering patterns of misogyny, negativity, and self-victimization in Incel discourse. Sentiment analysis, for example, has revealed widespread hostility in these spaces [13], while topic modeling has uncovered common themes and narratives that fuel discussions [2].

One major area of research in Incel communities is the identification of violent content. Studies like those by [21] have specifically examined toxic language across major Incel forums. Their analysis, which used a fine-tuned BERT model on 20,000 samples from toxic language datasets, found that expressions of hatred toward women were the most common form of toxicity. [21] also identified seven primary targets of hate, including women, society, and ethnic groups [21]. Broader studies, such as those by [2], have analyzed Incel ideologies in connection with acts of politically motivated violence. These studies show that some Incel communities encourage self-destructive behaviors and may inspire violent actions against women, successful men, or societal symbols representing perceived inequality. Further demographic analyses have explored the characteristics of Incel users, identifying common traits such as social isolation and engagement with mental health and relationship-related content [29].

However, many of these studies focus on smaller, isolated subsets of Incel users, missing a comprehensive understanding of behavior across larger forums. Research has shown that hateful content spreads more rapidly in these communities, which is crucial for understanding the broader dynamics of Incel communication [18]. Lundmark et al. explore the prevalence and characteristics of violent, threatful communication within online incel communities, with a particular focus on the Blackpill forum [16]. This research uses a machine learning model to analyze violent language within incel posts, focusing on three core aspects: the identity of perpetrators, the types of violence described, and the targets of violence [16]. The findings suggest that generalized perpetrators most commonly target entire groups, reflecting the dehumanization and ideological segregation inherent in these communities. Additionally, a significant portion of posts use coded language, such as references to video games, to obscure violent intentions, underscoring the challenge in detecting such speech in digital spaces. The paper highlights the complexity of identifying and addressing violent rhetoric due to these veiled threats and glorification of violence.

2.3 Blackpill.club: An Online Hub for Blackpill Ideology

Blackpill.club is an anonymous online forum that serves as a gathering space for individuals who adhere to the Blackpill ideology. The platform operates like a traditional message

board, where users create threads, reply to posts, and engage in discussions primarily centered around hopelessness in dating, resentment toward women, and fatalistic beliefs about social status and attractiveness [3]. Blackpill.club is known for promoting a hopeless outlook, where some members express the belief that trying to improve one's situation is pointless. However, this view is not an official position of the forum itself, but rather reflects the opinions of certain users communicating there.

According to a project by the Global Network on Extremism and Technology (GNET), platforms such as Blackpill.club, Incels.is, and Looksmax.org are moderated by the incels themselves and often contain extreme language, including misogyny, racism, homophobia, and the glorification of violence [12]. Discussions within Blackpill.club often focus on blaming women for men's struggles in the dating world [12]. Although the site has been criticized for promoting toxic ideas, particularly misogyny and harmful rhetoric, it remains a place for individuals who feel isolated or rejected by mainstream society to gather and express their views [27].

2.4 Understanding Dangerous Speech: A Framework for Identifying and Assessing Harmful Communication

Dangerous speech, as defined by [4], refers to “any form of expression (such as speech, text, or images) that increases the likelihood that its audience will justify or engage in violence against members of another group” [25]. The Dangerous Speech Project [25] developed a framework to help assess and identify potentially harmful speech. This framework includes five essential elements: Message, Audience, Context, Speaker, and Medium. The first element, Message, explores how the “out-group” is portrayed, which may include tactics like dehumanisation or labelling the out-group as a perceived threat [25]. Furthermore, [25] suggests that while not all harmful speech leads directly to violence, it significantly contributes to a cultural climate where violent acts are more likely to be normalized or justified. They emphasize the need to carefully examine the combination of these factors, as even speech that may seem harmless on its own can be dangerous in the right context.

3

Method

To address the research problem and answer the research questions, this thesis adopts a document survey as its primary research strategy. Surveys are advantageous for their ability to cover a broad spectrum, facilitating the collection of extensive data and offering a comprehensive overview of the subject matter [14, 8]. Given the broad scope of the research questions, which aim to examine various aspects of violent speech in the blackpill.club, and the need to analyze a large number of posts to identify meaningful insights, the survey method is ideal .

An alternative strategy considered was a case study method. Case studies allow for an in-depth examination of a particular instance or a narrowly defined aspect of a phenomenon [8]. This strategy is beneficial when the research aims to explore specific contexts or unique cases in detail. However, the current study's objectives are broader, aiming to encompass a wide range of data without the distinct boundaries that a case study typically requires. Moreover, case studies often focus on understanding the 'how' and 'why' of a situation within a specific setting [8], which does not align with the overarching goals of this thesis. Therefore, while a case study could provide deep insights into a singular aspect, it would not adequately address the comprehensive scope necessary for this thesis.

3.1 Data Collection

Documentary research is chosen as the data collection method, as it allows for the systematic extraction and analysis of textual data from online forums. In this study the documents consist of posts from blackpill.club. According to [8], posts from social media and discussion forums can be treated as documents, making this method suitable for utilizing them as primary data sources. This approach enables a comprehensive examination of discussions related to Incels and violent or threatening communication.

An alternative approach could have been observation, where researchers immerse themselves in the forum environment to understand participant interactions and cultural nuances [8]. However, given the extensive data set analyzed in this study, observation would

have been impractical. Observational methods are better suited for smaller-scale studies focusing on participant behaviour in real-time, whereas documentary research allows for efficient processing and systematic analysis of large datasets [8].

3.2 Data analysis

This study employs a mixed-method approach to analyze violent and threatening communication within Blackpill.club. By combining content analysis with both quantitative and qualitative methods, the study was able to capture a comprehensive understanding of the data and ensure a more nuanced evaluation of the content in these forums. This method was chosen because it allows for systematic coding and categorization of text data, enabling the researcher to quantify patterns and trends that directly address the research questions [8, 14]. This approach is well suited for Blackpill.club as the nature of the posts is highly unstructured, with vast amounts of user-generated text that may contain nuanced and implicit forms of violence or misogyny. Qualitative analysis provides deeper insights and context, improving the understanding and interpretation of the data, while the quantitative aspect offers clarity through numerical data and statistical analysis, enabling a clearer understanding of trends and patterns.

Discourse analysis would have been more appropriate for this study if the research question had been framed around investigating the purpose or intent of toxic narratives and what these messages aim to achieve for those who post them [8]. Since this study focuses on identifying and analyzing the frequency and severity of violent and threatening communication, content analysis was chosen because it better aligns with the objective of this study.

3.3 Ethical Considerations

As this thesis involves publicly available data from an online forum, ethical considerations remain essential, even if formal informed consent from individual users is not always feasible. According to Townsend and Wallace, researchers should assess whether users expect privacy and whether the data's nature requires additional ethical safeguards, even when the content is publicly accessible [32]. In line with best practices in research ethics, the following considerations will be prioritized:

- **Anonymity and confidentiality:** All posts analyzed will be anonymized to prevent identification of individual users. The study will remove any personal identifying information from the data to ensure the privacy of forum members.
- **Sensitive content:** Given the nature of the research, which involves violent or harmful language, care will be taken not to amplify or propagate any dangerous content. Posts that explicitly promote violence or harm will be flagged as part of the analysis, but the findings will focus on understanding the prevalence and nature of such communications rather than sharing them.
- **Avoiding harm:** The study will ensure that no posts are shared outside of the research context in a way that might encourage or disseminate harmful ideologies. The analysis will be limited to internal use and academic purposes only, ensuring that harmful rhetoric does not spread beyond the boundaries of the study.

- Respect for the community: While the study aims to understand how violence and threats manifest within the Blackpoll forum; this will avoid sensationalizing the content. The study will approach the topic with care and integrity, focusing on the academic and societal implications of the communication patterns within incel communities.

4

Application Of Methodology

4.1 Data Collection

The data set for this study consists of the post from blackpill.club. The original dataset, which was stored in a compressed JSON.zst file, contained more than 700,000 posts from the forum. For the purposes of this research, all comments posted between 2021 and 2024 amounting to approximately 350,000 posts were selected for detailed analysis.

4.1.1 Data Extraction and Field Selection

The original JSON.zst file included a variety of fields for each post. However, to streamline the analysis and ensure relevance to the research objectives, only the following fields were extracted:

- Post ID: A unique identifier for each post.
- Date and Time: The timestamp indicating when the post was created.
- Content: The main text or comment of the post.

These fields were selected based on the research focus on the temporal aspects and textual content of the posts, which were essential for analyzing patterns in violent and non-violent communication.

4.1.2 Data Processing Workflow

The data processing followed a structured approach to prepare the data set for analysis. The compressed.json.zst file was first decompressed using the appropriate tools. A Python script was then used to extract three relevant fields: post-identity, date and time, and content, from the JSON data. These fields were selected to align with the study focus on the timing and content of forum posts. The extracted data was subsequently converted into an Excel spreadsheet to facilitate further processing, filtering, and annotation. This process ensured that the data set was manageable and aligned with the study objectives.

4.1.3 Sampling Strategy

The sampling process for this study was carried out in two stages to ensure that the data set was representative of violent and threatening communication. Initially, a random sampling approach was employed to select 600 comments from a total of 350,000 comments. These comments were manually annotated to identify violent content and determine the frequency of violent speech, providing a reliable estimate of its prevalence within the dataset.

However, the random sampling approach did not adequately capture the volume of violent and threatening content necessary to analyze specific targets and severity levels. Therefore, a more targeted, keyword-based sampling strategy was adopted. Using a Python script, comments containing keywords such as ‘kill’, ‘murder’, ‘harm’, and ‘attack’—selected based on prior research and domain knowledge—were filtered. This process yielded a combined set of 12,445 comments (including the initial random sample).

To ensure the final dataset was statistically representative of this subset, the required sample size was recalculated using a 95% confidence level and a 3% margin of error, resulting in a sample size of 1,181 comments. These comments were then randomly selected from the filtered subset for further analysis, ensuring both statistical robustness and a focus on relevant violent and threatening content. Table 4.1 shows the stages of the sampling process used in this study.

Stage	Description	Number of Comments
Initial data collection	All comments collected from blackpill.club	700,000
Selected timeframe (2021–2024)	All comments posted between 2021 and 2024	350,000
Random sampling	Randomly selected to estimate how much violent communication is present	600
Filter for violent/threatening	Comments containing violent or threatening keywords	12,445
Sample size recalculation	Sample size calculated for analysis (95% confidence, 3% margin of error)	1,181
Final annotated dataset	Comments annotated manually and using AI	1,181

Table 4.1: Stages of the sampling and annotation process.

4.2 Data Analysis

Annotation of the dataset was conducted in two stages. In the first stage, the random sample of 600 comments was manually annotated to determine the presence of violent or threatening communication, providing a reliable estimate of its prevalence across the larger dataset. In the second stage, the keyword-based sample was used to analyze the severity and target groups of violent communication, as the larger pool of violent comments allowed for a more meaningful categorization. The researcher manually annotated the data and also used two

AI models, ChatGPT- 4.0 and Grok-3, to support the process. Multiple rounds of annotation were performed using varied prompts and batch sizes to improve consistency and accuracy. Each comment was annotated for three key variables: violence classification, severity level, and target group. To assess agreement between the annotators (the researcher and the two AI models), Krippendorff's alpha was calculated using JASP statistical software.

This two-stage analysis enabled the study to quantify the prevalence of violent communication, evaluate its intensity, and identify the groups most frequently targeted.

4.2.1 Frequency of Violent Comments

The frequency of violent speech was analyzed using the 600-comment subset drawn from the larger dataset to determine how common violent communication is within the forum. In this study, frequency refers to the proportion of comments that contain violent or threatening content relative to the total number of comments in the sample. Measuring frequency is important because it indicates the prevalence of harmful language in the forum and helps assess its potential impact. The aim of this analysis was therefore to identify how many comments were classified as violent compared with non-violent, and to evaluate whether violent speech constitutes a significant concern in the dataset.

Violence was defined according to the criteria adapted from [19], in which a comment is considered violent if it discusses harming others, encourages or praises violence, imagines violent acts, or threatens harm to others or oneself. This definition ensured consistency in the annotation process. The initial AI prompts used for annotation were developed based on this definition, and additional prompts tested during the annotation stages are provided in the Appendix.

4.2.2 Target Group Distribution

The target group distribution analysis focused on identifying which groups were most frequently targeted by violent comments. This analysis was conducted on the keyword-based subset of violent comments, which provided a sufficiently large sample for meaningful categorization. Target groups categorized as politicians, women, immigrants, Jews, non-white people, LGBTQ+ individuals, Muslims, and other, where the 'other' category includes all targets that do not fall within these predefined groups. This analysis was chosen because identifying the targets of violent speech is crucial to understanding who is most vulnerable to harm. By recognizing patterns in the targeting of specific groups, this analysis helps reveal potential biases or discriminatory attitudes expressed by the commenters within the dataset and how certain communities may be more at risk of being targeted by violent language.

4.2.3 Severity of Violent Comments

The severity of violent comments was analyzed to understand how strong or harmful the violent language in the dataset was. In this study, severity refers to how serious and direct a violent comment is. To measure this, each violent comment in the keyword-based sample was placed into one of three categories: low, moderate, or high severity. These categories were made using clear rules, making it easy to understand what type of comment fits each level.

Low severity includes comments that use violent or aggressive language but do not express a clear threat toward a person or group.

Moderate severity includes comments that call for harm or suggest that someone should be hurt, but do not describe a specific plan or direct threat.

High severity includes comments that make explicit, direct, or detailed threats of violence toward a person or group.

This classification helps show the difference between general hostile language and serious, threatening statements. Table 4.2 presents the three severity levels along with short descriptions and examples, making it easier for readers to understand how comments were categorized.

Severity Level	Description	Example
Low Severity	General violent language, no specific threat.	"I hate them all."
Moderate Severity	Calls for harm but no specific plan.	"Someone should beat them up."
High Severity	Explicit, detailed threat of harm.	"I have a gun, and I will shoot them tomorrow."

Table 4.2: Severity Levels of Violent Comments

5

Results

The findings presented in this chapter and the next include descriptions of violent and threatening rhetoric. To protect readers and maintain the privacy of the authors, no direct quotations from the analyzed posts are included.

5.1 Frequency of Violent Speech

From the initial random sample of 600 comments, manual annotation identified 52 comments (9%) as containing violent speech, while the remaining 548 comments (91%) were classified as non-violent. The keyword-based sampling method identified a total of 12,445 comments as potentially violent. Manual annotation of a random sample consisting of 1,181 comments from this set revealed that only 72% actually contained violent speech while 28% were classified as non-violent.




Figure 5.1: Distribution of violent and non-violent comments in the random sample of 600 comments.

5.2 Primary Targets of Violent speech

The analysis of targets of violent speech was conducted on the sample of 1,181 comments. These comments were selected from the larger dataset using a keyword-based sampling method and were manually annotated to identify specific targets and categorize the severity of violent speech. Out of the 1,181 comments, 851 were identified as violent. During the content analysis phase, these violent comments were further examined to determine the primary targets of the violent speech. The targets were categorized into eight groups: Women, Politicians, LGBTQ+ individuals, Immigrants, Jews, Black people, Muslims, and an “Other” category. The “Other” group includes various less frequently targeted groups such as incels, refugees, and government servants. Detailed descriptions of each target group can be found in Appendix A. Figure 5.2 presents a visual distribution of the percentage of comments targeting each group.

An analysis of the 851 violent comments on blackpill.club revealed that women were the most frequently targeted group, a total of 294 comments had women as the primary target. This was followed closely by a diverse “Other” category, which comprised about 40% (340 comments). Politicians were also targeted, though to a lesser extent, with around 7% (60 comments). Other groups such as LGBTQ+ individuals (5%, 45 comments), Non-white people (4%, 38 comments), Jews (4%, 36 comments), immigrants (3%, 23 comments), and Muslims (2%, 15 comments) appeared less frequently as targets.




Figure 5.2: Primary targets of violent speech

5.3 Severity Levels of Violent speech

To assess the intensity of violent language, a total of 1,181 comments selected by keyword-based sampling were annotated. Of these, 851 were identified as violent and subsequently categorized into one of three severity levels: high, moderate, or low. As shown in Figure 5.3, a substantial portion of the comments fell into the high severity category with almost 37% (312 comments). These comments often included very clear and specific threats, detailed descriptions of violence, or direct encouragement to hurt someone.

The moderate severity category accounted for 29% (252 comments). These comments typically expressed a desire to harm someone but in a more vague or indirect way, for example, suggesting that someone “should be hurt” without stating how or by whom. Finally, 33% (287 comments) were classified as low severity. These often included general expressions of hatred or hostility, such as saying they “hate” a certain group, without any specific threats or calls to action.




Figure 5.3: Severity Levels of Violent Comments (n=851)

5.4 Comparing ChatGPT and Grok in Annotating Violent Speech

To gain an understanding of how well ChatGPT and Grok can detect and annotate violent speech, their performance was evaluated using comments from Blackpill.club. A total of 1,181 comments were annotated by human annotation, with 851 identified as violent and further categorized by severity. Both AI tools were tasked with annotating these comments, and their results were compared to human annotations to evaluate accuracy and consistency. The overall inter-annotator agreement was calculated using Krippendorff’s alpha based on all annotated labels across the 1,181 comments. Table 5.1 presents the results

Comparison	Krippendorffs alpha
Grok vs. ChatGPT	0.602
Human vs. ChatGPT	0.521
Human vs. Grok	0.379

Table 5.1: Krippendorffs alpha Agreement Scores Between Annotators and AI Tools

The inter annotator agreement between ChatGPT and Grok was 0.602, indicating a moderate agreement. The agreement score between human annotators and ChatGPT was 0.521, while the agreement score between human annotators and Grok was lower, 0.379.

These results highlight that ChatGPT aligns more closely with human judgments in annotating violent speech than Grok. In practical terms, ChatGPT produced more relevant and context-sensitive annotations that directly reflected the input data. In contrast, Grok sometimes struggled to generate outputs that corresponded accurately to the input, occasionally producing annotations based on its own dataset rather than the comments provided, which likely contributed to its lower agreement scores. To further analyze annotation performance, a smaller batch size of 25 comments per prompt was used. The results of this evaluation are presented in Table 5.2.

Comparison	Krippendorffs alpha
ChatGPT vs. Human (small subset)	0.708
Grok vs. Human (small subset)	0.512
ChatGPT vs. Grok (small subset)	0.549

Table 5.2: Inter-annotator Agreement (Krippendorffs alpha) for batch size 25

When annotating small batches of 25 comments at a time using Excel files, both Grok and ChatGPT demonstrated improved performance compared to larger batches or single comment inputs. Table 5.3 shows the Krippendorffs alpha coefficients calculated on this subset and it revealed that ChatGPT exhibited strong agreement with human annotators at 0.708 when annotating comments as small batches. Grok's agreement with humans was moderate at 0.512, while the agreement between ChatGPT and Grok was 0.549, indicating moderate consistency between the two AI tools. Notably, ChatGPT's performance showed a decrease in agreement scores when annotating larger batches compared to these smaller batches, suggesting that smaller batch sizes enhance its annotation accuracy.

6

Discussion

This study explored the prevalence, severity, and primary targets of violent speech on blackpill.club. A key objective was evaluating the annotation performance of two AI models, ChatGPT and Grok, in detecting and categorizing violent content. By comparing their outputs against manual human annotations, the study evaluated the consistency and accuracy of these AI tools. The rest of this chapter addresses each research question in turn. Each question is presented in its own section to ensure a clear and focused discussion.

6.1 RQ1: How frequently does violent or threatening speech appear on blackpill.club?

To answer the first research question, the study examined how frequently violent or threatening comments appear on blackpill.club. In the initial random sample of 600 comments, manual annotation identified 52 comments (9%) as containing violent speech, while 548 comments (91%) were non-violent. This indicates that violent content, though present, represents a relatively small proportion of overall forum posts. Using the keyword-based sampling method, a total of 12,445 comments were flagged as potentially violent. Manual annotation of a random subset of 1,181 comments from this set revealed that 72% of these actually contained violent speech, while 28% were non-violent. The analysis shows that violent and threatening communication is a persistent feature of blackpill.club[22]. In the randomly sampled dataset, a considerable proportion of comments contained violent or threatening expressions, indicating that such language is not an isolated or occasional phenomenon[4]. The keyword-based sample also demonstrated a high concentration of explicit and implicit violent rhetoric.

The results also highlight the limitations of keyword-based sampling for estimating actual prevalence: while it captures a high density of violent content, it over-represents extreme expressions. The random sample provides a more accurate picture, demonstrating that although not every post is violent, the recurrent appearance of violent rhetoric shapes the communicative climate of the forum.

6.2 RQ2: Who are the primary targets of the violent speech?

The results clearly show that women are the most frequent targets of violent and threatening speech on blackpill.club. To address the research question on who is targeted by violent speech, the manually annotated comments were analyzed for the primary groups mentioned. The content analysis showed that women were the most frequently targeted group comprising 294 of the 851 comments annotated as violent, followed by a broad 'Other' category that included groups such as incels, refugees and government servants. This is consistent with prior research in incel forums, where misogyny is often a central ideological component [9]. Politicians, LGBTQ+ individuals, black people, Jews, immigrants, and Muslims were also identified as frequent targets, illustrating how violent rhetoric is spread across multiple marginalized groups.

The fact that women were the main targets in violent speech aligns with patterns found in some online male dominated communities, such as incel forums and blackpill.club[16]. In these groups, women are seen as the ones controlling access to sex and are portrayed as oppressors who destroy male identity [19][16]. This kind of language often starts with anger, but can escalate to dehumanizing women, suggesting that they deserve punishment or even elimination. Even though women were targeted the most, violent speech was also directed at politicians, immigrants, ethnic minorities, and LGBTQ+ people. This shows that hateful messages on these forums are not only about gender but also mix with racism, xenophobia, and anti LGBTQ+ ideas, similar to what is found in far-right online spaces. This targeting pattern reflects the deeply gendered and diverse nature of hostility in blackpill.club[24].

6.3 RQ3: How severe is the violent speech?

The analysis of the severity of violent comments provides insight into the intensity of aggressive and harmful content on blackpill.club. A significant proportion of the violent comments in the dataset fell into the high-severity category, indicating that explicit calls for or depictions of violence are relatively common. Among the manually annotated violent comments, over one-third (36%) were classified as high severity, containing explicit threats or direct encouragements of violence. Moderate severity comments included calls for harm without a specific plan, while low severity comments expressed general hostility without direct threats. These findings indicate that violent content on the forum is not only present but often intense, suggesting a highly toxic online environment. The high proportion of severe comments raises concerns about the normalization of aggressive rhetoric, which may contribute to the radicalization of users or increase the risk of real world harm. By distinguishing between low, moderate, and high severity, the study highlights the range of violent speech present, from general expressions of anger to explicit threats. This helps to answer the research question regarding the intensity of violent communication and emphasizes the potential danger posed by extreme hostility in online incel communities.

6.4 RQ4: How do ChatGPT and Grok differ in annotating violent speech in blackpill.club, and which tool is more accurate and efficient

To answer the fourth research question, this study compared the annotation performance of ChatGPT and Grok in identifying violent speech on blackpill.club. The comparison between ChatGPT and Grok shows clear differences in accuracy and reliability. ChatGPT demonstrated stronger agreement with human annotation, particularly in the smaller annotation batches, indicating a higher level of consistency and precision. Grok, in contrast, produced more variation and lower overall agreement scores in relation to both human annotators and ChatGPT.

In terms of AI annotation performance, ChatGPT demonstrated moderate to strong alignment with human annotators, achieving a Krippendorff's alpha of 0.521 overall and an improved 0.708 for smaller batch annotations. Grok, on the contrary, showed lower levels of agreement (0.379 overall and 0.512 in smaller batches), suggesting challenges in accurately capturing the nuanced language and context of violent speech within this dataset. These results emphasize ChatGPT's superior capacity to generate context aware annotations that closely reflect human judgment, making it a promising tool for content moderation and research applications in identifying harmful language online. However, the drop in agreement scores when ChatGPT annotated larger batches suggests that input size and format may significantly impact performance, underscoring the need for optimized annotation workflows. However, Grok's tendency to produce less relevant annotations when relying occasionally on patterns from unrelated datasets raises concerns about its reliability for sensitive tasks requiring precise contextual understanding. The moderate agreement (0.602) between ChatGPT and Grok also indicates some consistency, but reveals that the two AI models interpret the data differently, likely due to differences in training data and model architectures. The higher alignment of ChatGPT with human judgments suggests that it is better suited for tasks requiring a nuanced understanding of violent and hateful rhetoric. In contrast, Grok's struggles, including generating annotations influenced by unrelated datasets, highlight limitations that must be addressed before it can be reliably deployed for this kind of sensitive content analysis.

Efficiency was also considered and both models showed rapid processing capabilities. However, accuracy is the determining factor for the applicability in the real-world. Based on the results of this study, ChatGPT is the more reliable option for violence detection tasks, although neither AI model performs at a level that would completely replace human annotators. These findings highlight the importance of evaluating AI tools before deploying them in sensitive content moderation scenarios[26].

6.5 Limitations

Using automated tools like ChatGPT and Grok to detect hate or violent speech is not always reliable. These tools can sometimes miss harmful content or label harmless posts as toxic. Because of this, manual analysis is often needed to check and correct these mistakes. However, manual annotation also has its own challenges. Understanding hate or violent speech

can be subjective what one person sees as offensive, another person might not. Even if annotators follow the same rules, their personal views or experiences can affect how they label the posts.

In this study, only one person (the researcher) did the manual annotation. While care was taken to follow clear guidelines, it is still possible that another researcher could label the same posts differently or focus on other aspects of the content. Also, people may describe similar groups using different terms (for example, saying “immigrants” instead of a specific ethnic group), which can change how a post is interpreted. Because of these issues, it is unlikely that someone else doing the same study with different prompts or labeling choices would get exactly the same results. Therefore, the findings should be seen as an informed interpretation, rather than a precise or final measurement of toxic content.

6.6 Social and Ethical Implications

Beyond research ethics, this study highlights important societal and ethical issues. Violent and threatening communication on blackpill.club, although not present in most posts occurs frequently enough to target women and other marginalized groups. Understanding these narratives is important because exposure to repeated hostility can normalize misogyny, reinforce harmful stereotypes, and increase risks of online harassment or offline radicalization. The study should not be seen as an effort to restrict freedom of speech. Rather, it aims to understand the content and severity of violent or threatful posts, and the groups affected by such narratives. Care was taken to ensure that presenting these findings does not provide a platform for spreading these violent ideas[8]. The use of AI tools like ChatGPT and Grok to detect and classify violent speech also raises ethical considerations. These models can misclassify nuanced or coded language, reflect biases in training data, or disproportionately flag certain opinions as violent. Platforms and researchers must therefore combine AI with human oversight, ensure transparency in moderation processes, and carefully weigh the potential societal impact of interventions[31]. Overall, this study shows that research into online discourse can contribute to a better understanding of how violent and threatening language functions in online communities, while highlighting the need for careful and responsible use of AI tools when addressing such content.

6.7 Conclusion

This study set out to explore the nature of violent speech on blackpill.club, focusing on its frequency, targets, severity, and how well AI tools can detect such content. The findings show that violent speech does not dominate the forum but still accounts for a substantial and concerning share of posts (around 9%), highlighting the need to monitor and address harmful expressions within the community.

In terms of targets, women were the primary focus of violent speech, followed by politicians, immigrants, LGBTQ+ individuals, and groups such as Black people and Jews. This pattern indicates a strong presence of misogynistic and xenophobic narratives. Regarding severity, a significant portion of violent comments included highly toxic language, such as dehumanizing metaphors and explicit calls for violence, emphasizing the intensity and potential harm of the rhetoric used. When evaluating AI tools, both ChatGPT and Grok were

able to identify and classify violent speech efficiently, but their accuracy varied. ChatGPT demonstrated higher consistency and closer alignment with human annotation than Grok. However, neither model achieved a level of accuracy sufficient to fully replace human judgment, particularly when handling nuanced or coded language[19].

Overall, the study confirms that violent and threatening communication is a recurring feature of blackpill.club, primarily targeting women but also other social and political groups, and often expressed in highly aggressive or explicitly violent ways. These findings underscore the need for detection strategies that combine AI tools with human oversight, especially in online spaces where language is complex, ironic, or coded. Continued investigation into harmful online communities is essential to understand how their discourse shapes social attitudes and potential real-world risks.

6.7.1 Future Research

In addition to studying violent speech, future research should also explore the presence and impact of self harm and suicidal language on blackpill.club. The findings indicated a notable amount of content expressing self-harm or suicidal thoughts, which suggests an urgent need to better understand the mental health risks within these communities.

Expanding the scope to investigate the severity levels of such content more broadly would provide valuable insights into how often and in what ways users express distress or harmful intentions. Furthermore, analyzing demographic information such as age groups could help identify vulnerable populations and tailor intervention or support strategies more effectively.

As AI technology continues to improve, it is essential to leverage its capabilities for detecting harmful content online. Future research could explore different AI models or develop more advanced approaches to automatically identify toxic, violent, and self harm related speech. While manual annotation provides more accurate results, it is time consuming and challenging to apply to large datasets. Using AI tools more effectively could help overcome these limitations and support scalable content moderation and research efforts.

By combining these approaches, future studies could offer a more comprehensive view of both the violent and self-destructive narratives present on Blackpill.club, which could help identify patterns of online behavior linked to psychological distress and inform mental health awareness campaigns or intervention strategies

6.8 Relevance to Cybersecurity

This research contributes to the field of cybersecurity by advancing the understanding of how violent and hateful speech manifests in online extremist communities. By combining manual annotation with AI tools like ChatGPT and Grok, the study demonstrates effective methods for detecting harmful content in large datasets, which is crucial for developing automated moderation systems. The insights into the targets, severity, and patterns of violent language help inform better strategies for monitoring and mitigating online threats. Overall, the findings support the creation of more robust cybersecurity frameworks aimed at reducing online radicalization and protecting vulnerable groups from digital harm.

6.9 Use of AI tools

Throughout this thesis, various AI tools were used to assist with research and writing. Grammarly helped correct grammar and spelling, enhancing the clarity of the text. For data analysis, ChatGPT and Grok were used to annotate comments by detecting and categorizing violent speech. These AI models reviewed large volumes of comments, labeling them based on the presence and severity of violent language, which sped up the process compared to manual annotation alone. DeepSeek supported the research by locating relevant studies and background information. ChatGPT also helped improve the writing style without influencing the study's findings.

Bibliography

- [1] Anti-Defamation League. *The Extremist Medicine Cabinet: A Guide to Online “Pills”*. Nov. 2019. URL: <https://www.adl.org/blog/the-extremist-medicine-cabinet-a-guide-toonline-pills>.
- [2] S. J. Baele, L. Brace, and T. G. Coan. “From “Incel” to “Saint”: analyzing the violent worldview behind the 2018 Toronto attack”. In: *Terrorism and Political Violence* 33 (2021), pp. 1667–1691. DOI: 10.1080/09546553.2019.1638256.
- [3] Zoe Beauchamp. *Our incel problem: How a support group for the dateless became one of the internet’s most dangerous subcultures*. Apr. 2019. URL: <https://www.vox.com/thehighlight/2019/4/16/18287446/incel-definition-redd>.
- [4] S. Benesch. *Dangerous Speech: A Proposal to Prevent Group Violence*. <https://dangerousspeech.org>. Accessed: 2018-10-10. 2013. perma.cc: X8DP-M7Q2.
- [5] Anthony G. Burton. “Blackpill Science: Involuntary Celibacy, Rational Technique, and Economic Existence under Neoliberalism”. In: *Canadian Journal of Communication* 47.4 (2022), pp. 676–701. DOI: 10.3138/cjc.2022-07-25. URL: <https://doi.org/10.3138/cjc.2022-07-25>.
- [6] G. L. Casey. “Ending the Incel Rebellion: The Tragic Impacts of an Online Hate Group”. In: *Loyola Journal of Public Interest Law* 21 (2019), p. 71.
- [7] Moonshot CVE. *Incels: A Guide to Symbols and Terminology*. May 2020. URL: http://moonshotcve.com/wp-content/uploads/2020/05/Incels_-A-Guide-to-Symbols-and-Terminology_Moonshot-CVE.pdf.
- [8] Martyn Denscombe. *The Good Research Guide: For Small-Scale Social Research Projects*. 5th. Maidenhead: Open University Press / McGraw-Hill Education, 2014.
- [9] Mohammad Hajarian and Zahra Khanbabaloo. “Toward Stopping Incel Rebellion: Detecting Incels in Social Media Using Sentiment Analysis”. In: *Proceedings of the 7th International Conference on Web Research (ICWR)*. Leganes, Spain: IEEE, 2021, p. 169. DOI: 10.1109/ICWR51843.2021.00041.
- [10] Hope not Hate. “Superconspiracies: Qanon and the new world order”. In: *Antisemitism in the Digital Age: online antisemitic hate, holocaust denial, conspiracy ideologies and terrorism in Europe*. Ed. by Joe Mulhall. Amadeu Antonio Foundation, Expo Foundation and Hope not Hate, 2021.
- [11] Brian Hoffman, Jessica Ware, and Eric Shapiro. “Assessing the Threat of Incel Violence”. In: *Studies in Conflict & Terrorism* 43.7 (2020), pp. 565–587. DOI: 10.1080/1057610X.2020.1769382.

- [12] International Centre for the Study of Radicalisation (ICSR). *Title of the Report*. Tech. rep. King's College London: Global Network on Extremism and Technology (GNET), 2024. URL: <https://www.examplelink.com>.
- [13] S. Jaki et al. "Online hatred of women in the Incels.me forum: linguistic analysis and automatic detection". In: *Journal of Language Aggression and Conflict* 7 (2019), pp. 240–268. DOI: 10.1075/jlac.00026.jak.
- [14] P. Johannesson and E. Perjons. *An Introduction to Design Science*. Elektronisk resurs. Cham: Springer International Publishing, 2014.
- [15] D. Lazer et al. "Computational social science". In: *Science* 323 (2009), pp. 721–723. DOI: 10.1126/science.1167742.
- [16] Lukas Lundmark, Lisa Kaati, and Amendra Shrestha. "Visions of Violence: Threatful Communication in Incel Communities". In: *Proceedings of the 2024 IEEE International Conference on Big Data*. Stockholm, Sweden: IEEE, 2024. DOI: 10.1109/BigData62323.2024.10824947.
- [17] Amanda Marcotte. "How 'Pick-Up Artist' Philosophy and Its More Misogynist Backlash Shaped Mind of Alleged Killer Elliot Rodger". In: *The American Prospect* (May 2014). URL: <https://prospect.org/culture/pick-up-artist-philosophy-misogynist-backlash-shaped-mindalleged-killer-elliott-rodger/>.
- [18] B. Mathew et al. "Spread of hate speech in online social media". In: *Proceedings of the 10th ACM Conference on Web Science*. New York, NY: ACM, 2019, pp. 173–182. DOI: 10.1145/3292522.3326034.
- [19] Daniel Matter et al. "Investigating the increase of violent speech in Incel communities with human-guided GPT-4 prompt iteration". In: *Frontiers in Sociology and Psychology* 2 (2024). First authorship shared between Matter and Schirmer, p. 1383152. DOI: 10.3389/frsps.2024.1383152. URL: <https://doi.org/10.3389/frsps.2024.1383152>.
- [20] Kostantinos Papadamou et al. "Understanding the Incel Community on YouTube". In: *Proceedings of the ACM on Human-Computer Interaction* 5 (Oct. 2021), pp. 1–25. DOI: 10.1145/3479556.
- [21] B. Pelzer et al. "Toxic language in online incel communities". In: *SN Social Sciences* 1 (2021), pp. 1–22. DOI: 10.1007/s43545-021-00220-8.
- [22] Shruti Phadke, Mattia Samory, and Tanushree Mitra. "Pathways through Conspiracy: The Evolution of Conspiracy Radicalization through Engagement in Online Conspiracy Discussions". In: *Proceedings of the International AAAI Conference on Web and Social Media*. Vol. 16. 2022, pp. 770–781. DOI: 10.1609/icwsm.v16i1.19333.
- [23] Kayla Preston, Michael Halpin, and Finlay Maguire. "The Black Pill: New Technology and the Male Supremacy of Involuntarily Celibate Men". In: *Men and Masculinities* 24.5 (2021), pp. 823–841. DOI: 10.1177/1097184X211017954. URL: <https://journals.sagepub.com/home/jmm>.
- [24] Lin Prøitz, Fredrik Langeland, and Henry Mainsah. *From Boyhood to Radicalised Masculinity: A Rapid Research Response from the International Symposium on Incels*. Nora Forlag, 2021.

- [25] Dangerous Speech Project. *Dangerous speech: a practical guide*. [Accessed: 2023-03-31]. 2021. URL: <https://dangerousspeech.org/guide/>.
- [26] Steven Rathje et al. *GPT is an Effective Tool for Multilingual Psychological Text Analysis*. PsyArXiv Preprint. 2023. DOI: 10.31234/osf.io/sekf5.
- [27] B. Rouda and A. Siegel. “‘I’d kill for a girl like that’: The Black Pill and the Incel Uprising”. In: *Journal of Online Communities* (2020).
- [28] Ian Steadman. “The Sexist Pseudoscience of Pick-Up Artists: The Dangers of ‘Alpha Male’ Thinking”. In: *The New Republic* (June 2014). URL: <https://newrepublic.com/article/118036/sexist-pseudoscience-alpha-male-pick-artists>.
- [29] S. Stijelja and B. L. Mishara. “Characteristics of Incel forum users: social network analysis and chronological posting patterns”. In: *Studies in Conflict and Terrorism* (2023), pp. 1–21. DOI: 10.1080/1057610X.2023.2208892.
- [30] Jim Taylor. “The woman who founded the ‘incel’ movement”. In: *BBC Radio 5 Live* (Aug. 2018). Accessed: 2025-03-22. URL: <https://www.bbc.co.uk/radio5live>.
- [31] David L. Townsend and Claire Wallace. *Social Media Research: A Guide to Ethics*. Aberdeen: University of Aberdeen, 2016.
- [32] Leanne Townsend and Claire Wallace. “The ethics of using social media data in research: A new framework”. In: *The Ethics of Online Research*. Emerald Publishing Limited, 2017.

Appendix A — Targeted Group

Target Group	Description
Politicians	Posts targeting individual politicians or political parties. Includes Swedish and international politicians (e.g., EU, US).
Women	Posts expressing hatred or toxic language directed at women or feminists. Does not include posts targeting female politicians or professionals.
Immigrants	Posts where the target is referred to as an “immigrant” without specifying ethnic or national origin.
Muslims	Posts referring to individuals as Muslims, often using derogatory or demeaning language.
Jews	Posts targeting Jewish people primarily as an ethnic group, not a religious one.
LGBTQ+	Posts directed at individuals identifying as lesbian, gay, bisexual, transgender, queer, etc. Most comments targeted gay or trans individuals.
Non-white people	Includes people of African, Middle Eastern, Asian, or South American descent. Posts often used racial slurs or focused on non-white physical appearance.
Other	Includes Men and incels, criminals, Romani people, Russians, Ukrainians, Religion, Europeans, Canadians and fat people

Appendix B — Python Code for Keyword Sampling

```
import pandas as pd
import re

try:
    df = pd.read_csv("extracted_data(Sheet1).csv", encoding='ISO-8859-1', delimiter=";",
                     on_bad_lines='skip')
except Exception as e:
    print("Error loading CSV:", e)
    exit()

df.columns = df.columns.str.strip().str.lower()

print("Available columns:", df.columns.tolist())

df['text'] = df['text'].fillna('').astype(str)
df['cleaned_text'] = df['text'].str.lower()

violent_keywords = ["kill", "murder", "torture", "destroy", "death", "stab", "shoot"]
threatful_keywords = ["harm", "attack", "danger", "threat", "target", "injure", "strike"]

def categorize_comment(text):
    def match_keywords(keywords):
        return any(re.search(r'\b' + re.escape(word) + r'\b', text) for word in keywords)

    if match_keywords(violent_keywords):
        return 'violent'
    elif match_keywords(threatful_keywords):
        return 'threatful'
    else:
        return 'uncategorized'

df['category'] = df['cleaned_text'].apply(categorize_comment)

violent_comments = df[df['category'] == 'violent']
threatful_comments = df[df['category'] == 'threatful']
```

```
violent_comments.to_csv('violent_comments.csv', index=False)
threatful_comments.to_csv('threatful_comments.csv', index=False)

print("\n--- Category Summary ---")
print(df['category'].value_counts())
```

Appendix C —Prompts Used for AI Annotation

Prompt 1 — Violent Speech Classification

You are an annotation assistant. Review each forum comment and classify whether it contains violent or threatening speech.

A comment is **violent** if it explicitly or implicitly threatens harm, encourages violence, or describes violent acts.

A comment is **non-violent** if it contains no threats or encouragement of harm.

Respond only with one label per comment: Violent or Non-violent

Prompt 2 — Severity Categorization

For comments labeled as violent, classify the severity of the violence.

- **High severity** – explicit threats, encouragement of killing, or detailed violent fantasies.
- **Moderate severity** – indirect threats or approval of harm without explicit calls to violence.
- **Low severity** – hostile or aggressive language without clear threats.

Respond with one label: High, Moderate, or Low.

Prompt 3 — Target Identification

Identify who or what the violence is directed toward.

Choose one or more of the following categories:

- Women
- Politicians
- Immigrants
- Jews
- Non-white people
- LGBTQ+ individuals
- Muslims

- Other / Unclear

Each batch was reviewed to ensure consistent formatting and interpretation before moving to the next stage.